

COMSM0007 Assignment: A Summary of Public Key Encryption with keyword Search

Dominic Moylett - dm1905@my.bristol.ac.uk

April 6, 2015

1 Introduction

Public Key Encryption with keyword Search is a paper by Dan Boneh, Giovanni Di Crescenzo, Rafail Ostrovsky and Giuseppe Persiano published at Eurocrypt 2004. This paper looks at the ability to search for items contained within encrypted data.

An example use is an email server: Alice can grant the server the ability to check if the email has the keyword 'Urgent'. When the server now receives an email, it checks if it has the keyword 'Urgent': if so the email is routed to Alice's pager, otherwise it is sent to her laptop. Non-interactive public key encryption with keyword search (sometimes abbreviated to 'searchable encryption') schemes mean that the server can check an encrypted email for keywords without learning what the keywords themselves are.

Due to the computational cost of public key encryption in practice, a searchable encryption scheme does not allow a computer to search the entire encrypted message for keywords. Instead, Alice gives the machine is given access to a list of keywords that they are able to search for. When Bob sends an email, he then selects a number of words to be keywords, and sends them using a *Public-Key Encryption with keyword Search* (PEKS) of each keyword. When the server receives the email, it checks if any of the keywords Alice has given it the ability to search for are within the list of keywords sent by Bob.

2 Definitions

Definition 1. *A Non-interactive public key encryption with keyword search scheme consists of the following polynomial time randomised algorithms:*

1. **KeyGen(s):** *Given a security parameter s , produce public and private keys A_{pub}, A_{priv} respectively.*

2. $\text{PEKS}(A_{\text{pub}}, W)$: Given a public key A_{pub} and keyword W , produce a searchable encryption of W .
3. $\text{Trapdoor}(A_{\text{priv}}, W)$: Given a private key A_{priv} and keyword W , produce a trapdoor T_W .
4. $\text{Test}(A_{\text{pub}}, S, T_W)$: Given a public key A_{pub} , searchable encryption $S = \text{PEKS}(A_{\text{pub}}, W')$ and trapdoor $T_W = \text{Trapdoor}(A_{\text{priv}}, W)$, output 'yes' if $W = W'$ and 'no' otherwise.

When Bob sends an email msg with keywords W_1, \dots, W_k to Alice, he computes the searchable encryptions for each keyword and sends the following:

$$[E_{A_{\text{pub}}}[\text{msg}], \text{PEKS}(A_{\text{pub}}, W_1), \dots, \text{PEKS}(A_{\text{pub}}, W_k)]$$

The email is received by Alice's server, which has trapdoors $T_{W'_1}, \dots, T_{W'_l}$ where $T_{W'} = \text{Trapdoor}(A_{\text{priv}}, W')$ for W' in keywords $[W'_1, \dots, W'_l]$. The server then uses *mathsf{Test}* with the trapdoors it has been given to see if the email contains one of the keywords specified by Alice.

Accompanying this definition is a security definition, specified as a PEKS Security game against an active attacker \mathcal{A} . The game is defined similarly to IND-CCA:

1. The challenger runs $\text{KeyGen}(s)$ to generate A_{pub} and A_{priv} . A_{pub} is given to the attacker.
2. The attacker can adaptively ask for the trapdoor T_W of any keyword $W \in \{0, 1\}^*$ of his choice.
3. At some point, the attacker \mathcal{A} sends two keywords W_0, W_1 . The only restriction is that the attacker must have not previously asked for the trapdoors T_{W_0} or T_{W_1} . The challenger picks a random bit $b \in \{0, 1\}$ and gives the attacker $C = \text{PEKS}(A_{\text{pub}}, W_b)$. C is the challenge PEKS.
4. The attacker can continue to ask for trapdoors T_W for any W as long as $W \neq W_0, W_1$.
5. The attacker eventually outputs a response bit $b' \in \{0, 1\}$ and wins if $b = b'$.

The adversary's advantage for winning the game is defined as:

$$\text{Adv}_{\mathcal{A}}(s) = |\Pr[b' = b] - \frac{1}{2}|$$

Definition 2. A PEKS is semantically secure against an adaptive chosen keyword attack if for any polynomial time attacker \mathcal{A} we have that $\text{Adv}_{\mathcal{A}}(s)$ is a negligible function.

3 Constructions

The remainder of the paper gives two constructions of public key searchable encryption schemes.

The first is based on bilinear maps. Bilinear maps are functions $e : G_1 \times G_1 \rightarrow G_2$ for groups G_1, G_2 of prime order p which are efficiently computable and satisfy the properties that for any integers $x, y \in [1, p]$ it holds that $e(g^x, g^y) = e(g, g)^{xy}$ and that if g is a generator of G_1 then $e(g, g)$ is a generator of G_2 . Using e and two hash functions $H_1 : \{0, 1\}^* \rightarrow G_1$ and $H_2 : G_2 \rightarrow \{0, 1\}^{\log p}$, a searchable encryption scheme can be devised.

The scheme is proven semantically secure against an adaptive chosen keyword attack in the random oracle model. This is proven via a reduction to the Bilinear Diffie-Hellman Problem: if an adversary \mathcal{A} can break the scheme with non-negligible advantage ϵ , an adversary \mathcal{B} can be constructed which, given $g, g^a, g^b, g^c \in G_1$ as input, can compute $e(g, g)^{abc} \in G_2$ with advantage $\frac{\epsilon}{eq_T q_{H_2}}$, where q_T and q_{H_2} are the number of calls to `Trapdoor` and H_2 respectively.

The second construction is based on any trapdoor permutation. A result from Bellare et al. states that given any trapdoor permutation, we can construct a source indistinguishable public key encryption scheme. A public key encryption scheme is source indistinguishable if, given a ciphertext, it is computationally difficult to determine the public key that encrypted it. The scheme works by A_{pub} being a set of public keys - one for each keyword - and A_{priv} being their corresponding private keys. The security of this scheme is shown by a reduction: If the keywords can be distinguished, the underlying encryption scheme is not source indistinguishable.

The advantage of this second scheme is that it does not rely on random oracles. But it has a cost in that the public and private keys grow linearly in size as the number of the keywords increase. This is improved by re-using source indistinguishable keys, which can still offer security via cover-free families. The construction uses parameters (d, t, k, q) where k is the number of words in the dictionary, t is an upper bound on the number of keyword trapdoors given to the server by Alice and d and q are two integers satisfying $q = \lceil \frac{d}{4t} \rceil$ and $d \leq 16t^2(1 + \frac{\log(k/2)}{\log 3})$ and is based on a lemma by Du and Hwang that there exists a deterministic algorithm that, for any (d, t, k, q) satisfying the above bounds, produces a q -uniform t -cover-free family. A reduction shows that if the PKES scheme has an adversary with advantage ϵ , then there exists an algorithm that breaks source indistinguishable public key encryption scheme and aborts with probability $\frac{1}{\text{poly}(t, q, d)}$. Running this algorithm repeatedly until it doesn't abort yields an expected polynomial time adversary with advantage $\frac{\epsilon}{q^2}$.

4 Context